


MECS 6616

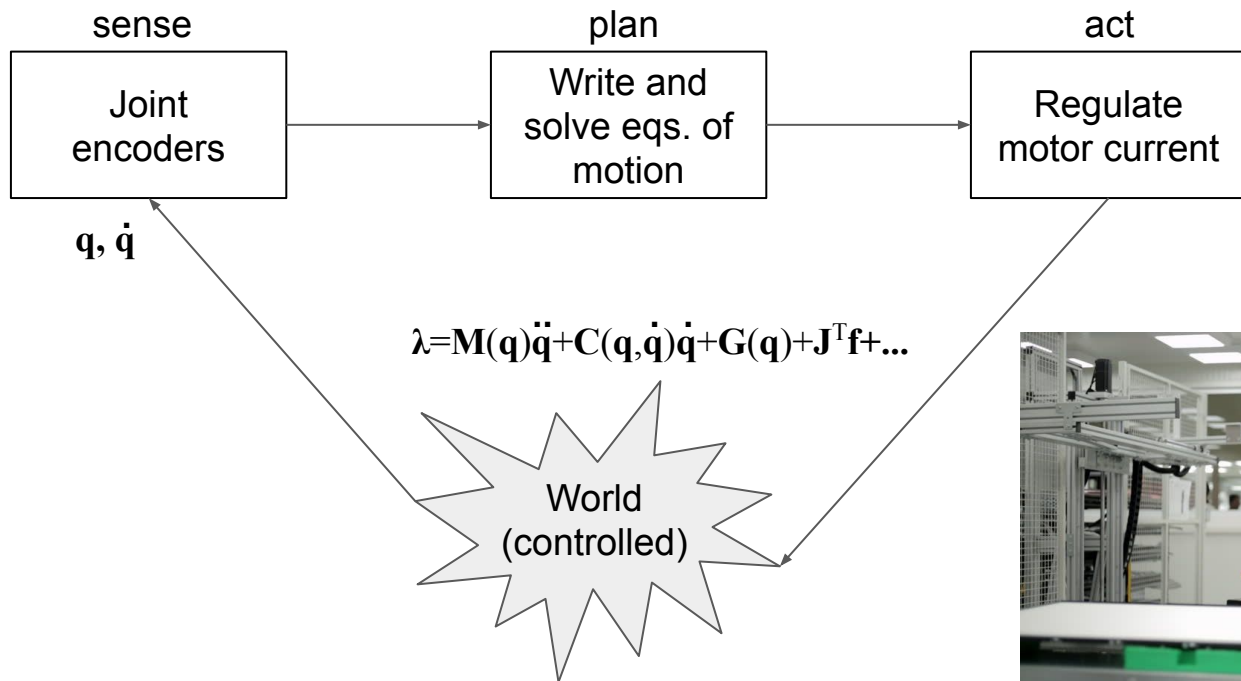
Robot Learning

Spring 2020
Matei Ciocarlie

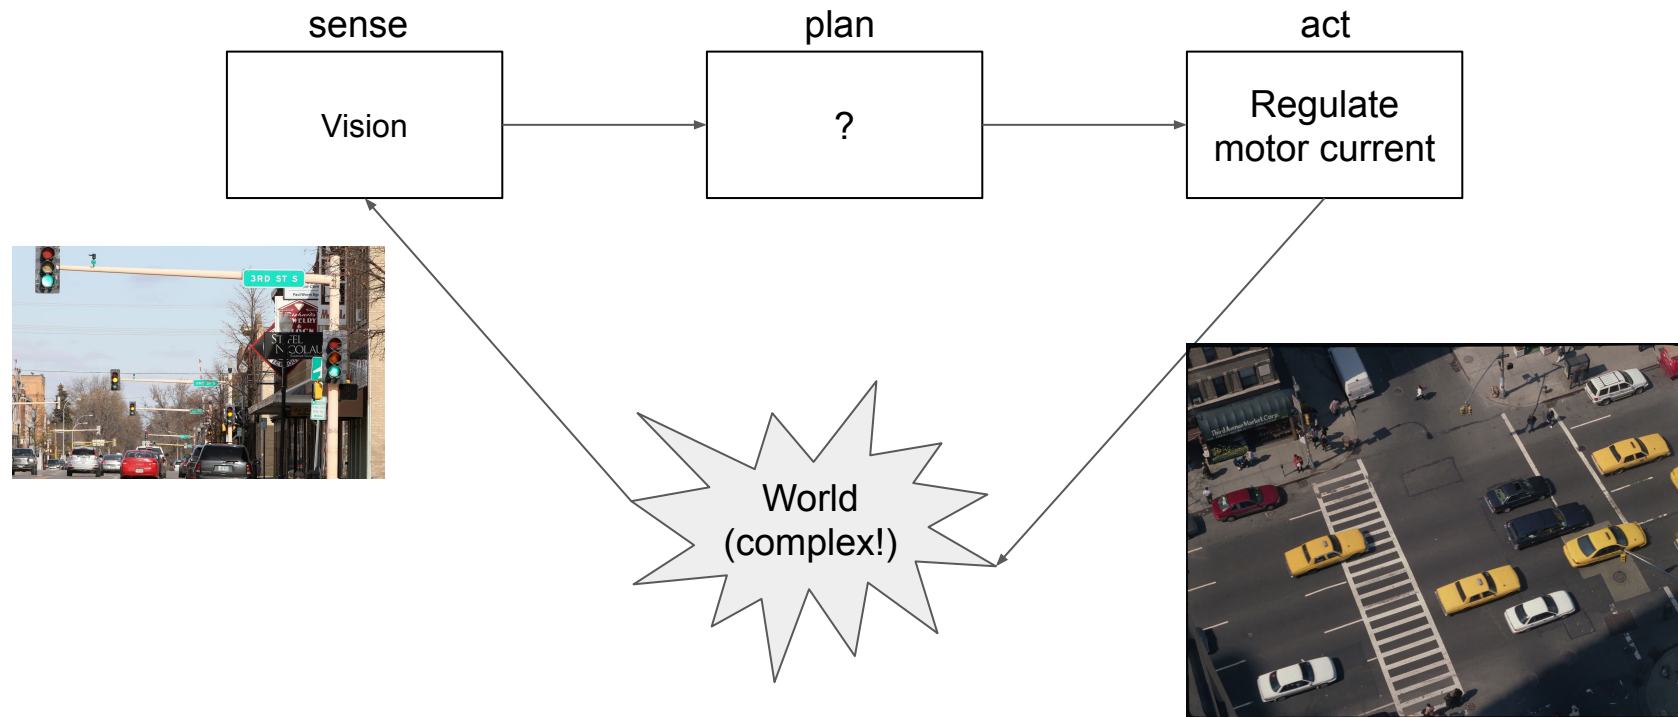
The Robotics Loop



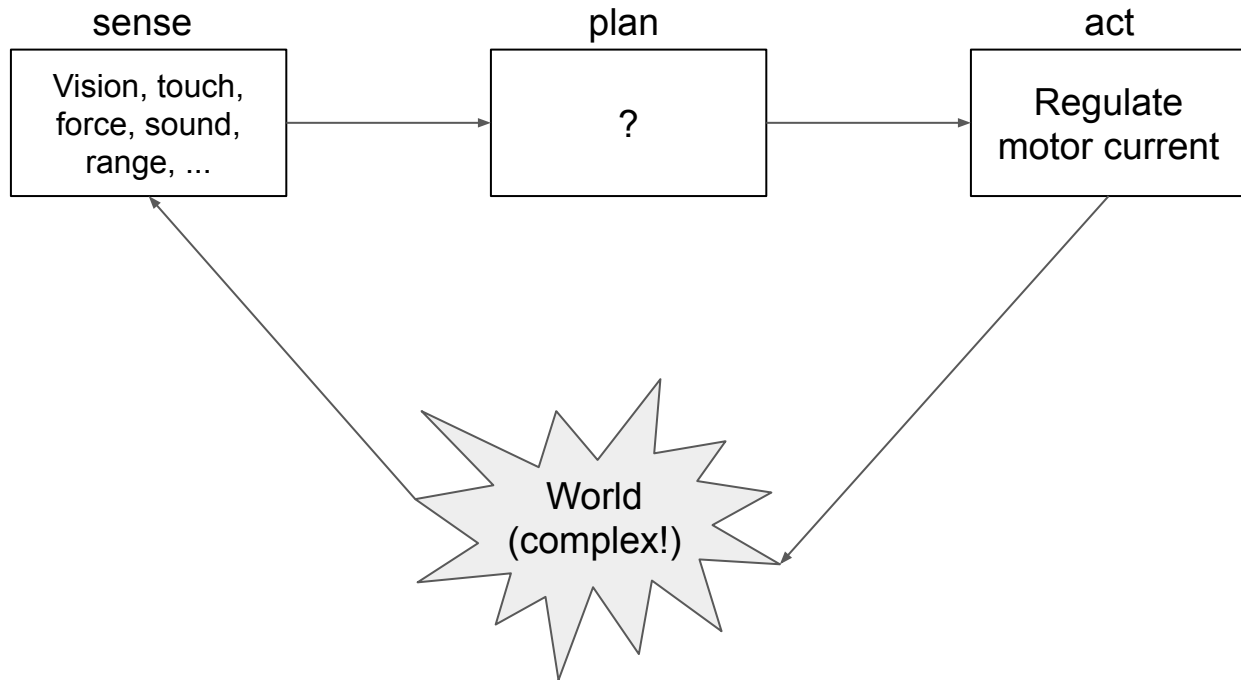
The Robotics Loop



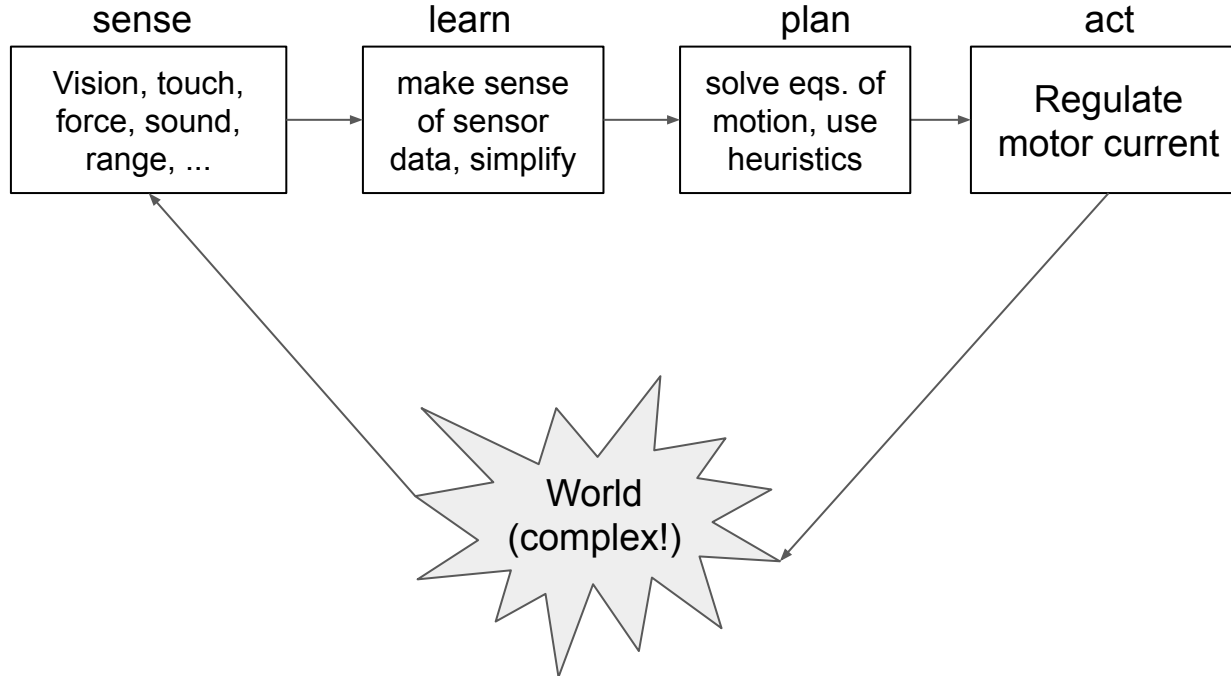
The Robotics Loop



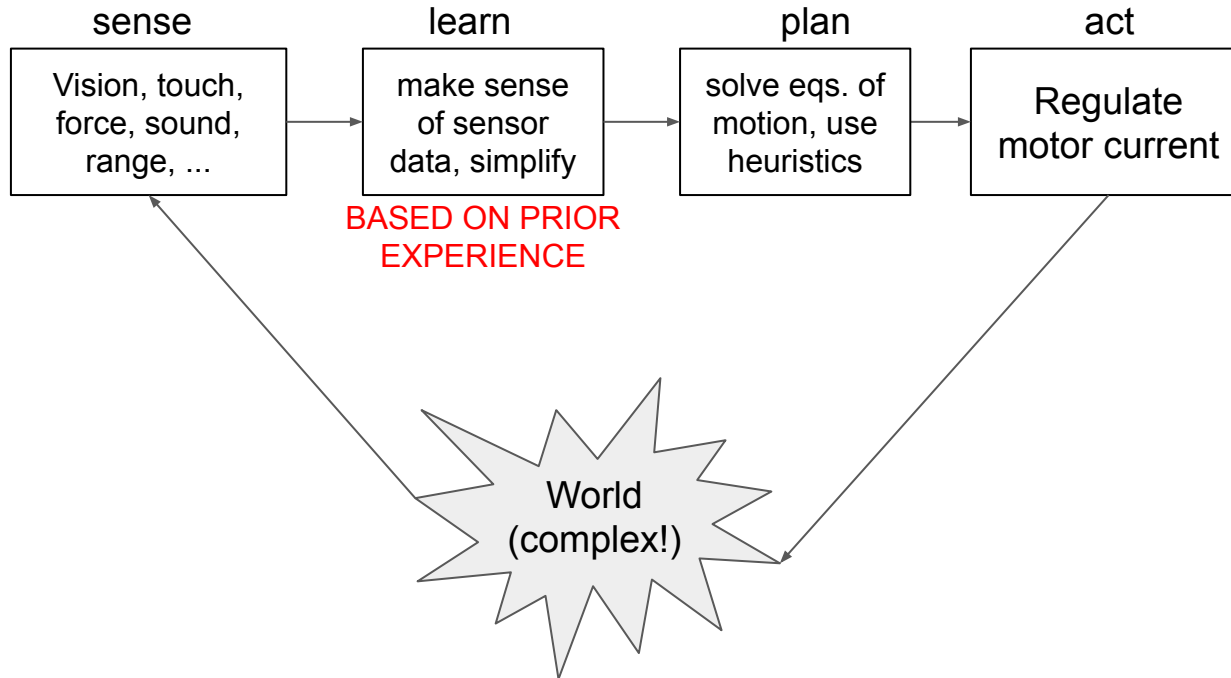
The Robotics Loop



The Robotics Loop



The Robotics Loop




What is Machine Learning?



Option A: magic


You put in your problem and out comes the answer, and you don't need to write down a single equation about your robot.



Option B: statistical estimation

You fit curves to points.


What is Machine Learning? Unfortunately:




Option B: statistical estimation

You fit curves to points.

Statistical Estimation




Statistical Estimation




Interpolation

Statistical Estimation




Interpolation

Statistical Estimation



Extrapolation

Statistical Estimation



Extrapolation

Broadest ML Problem Categorization

- Unsupervised
- Supervised

Fundamental Unsupervised ML Problems

$$\mathbf{x} \in \mathbb{R}^{d1}$$

Find some **structure** in the data you are given, **simplify**

$$\mathbf{x}_1$$

Clustering:

$$\mathbf{x}_2$$

Find a number of discrete classes, or clusters:

...

$\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k, k \ll n$

$$\mathbf{x}_n$$

Assign each point to a cluster

$\mathbf{x}_k \in \mathbf{c}_i$ (“belongs” to \mathbf{c}_i)

Minimize “distance” between points and their cluster

Fundamental Unsupervised ML Problems

$$\mathbf{x} \in \mathbb{R}^{d1}$$

Find some **structure** in the data you are given, **simplify**

$$\mathbf{x}_1$$

Clustering:

Dimensionality reduction:

$$\mathbf{x}_2$$

Find a number of discrete classes, or clusters:

Express each data point using fewer variables

...

$\mathbf{c}_1, \mathbf{c}_2, \dots, \mathbf{c}_k, k \ll n$

$$\mathbf{x}_n$$

Assign each point to a cluster

$\mathbf{x}_k \in \mathbf{c}_i$ (“belongs” to \mathbf{c}_i)

$$\mathbf{x}_i \sim f(\mathbf{y}_i), \mathbf{y}_i \in \mathbb{R}^{d2}, d2 < d1$$

Minimize information loss:

$$\|\mathbf{x}_i - f(\mathbf{y}_i)\|$$

Minimize “distance” between points and their cluster

Fundamental Supervised ML Problems

Fundamental Supervised ML Problems

$$\mathbf{x} \in \mathbb{R}^{d_1}, \mathbf{y} \in \mathbb{R}^{d_2}$$

$$\mathbf{x}_1 \rightarrow \mathbf{y}_1$$

$$\mathbf{x}_2 \rightarrow \mathbf{y}_2$$

...

$$\mathbf{x}_n \rightarrow \mathbf{y}_n$$

Fundamental Supervised ML Problems

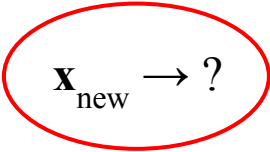
$$\mathbf{x} \in \mathbb{R}^{d_1}, \mathbf{y} \in \mathbb{R}^{d_2}$$

$$\mathbf{x}_1 \rightarrow \mathbf{y}_1$$

$$\mathbf{x}_2 \rightarrow \mathbf{y}_2$$

...

$$\mathbf{x}_n \rightarrow \mathbf{y}_n$$


$$\mathbf{x}_{\text{new}} \rightarrow ?$$

Fundamental Supervised ML Problems

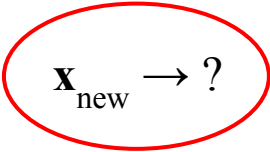
$$\mathbf{x} \in \mathbb{R}^{d_1}, \mathbf{y} \in \mathbb{R}^{d_2}$$

$$\mathbf{x}_1 \rightarrow \mathbf{y}_1$$

$$\mathbf{x}_2 \rightarrow \mathbf{y}_2$$

...

$$\mathbf{x}_n \rightarrow \mathbf{y}_n$$


$$\mathbf{x}_{\text{new}} \rightarrow ?$$

Classification:

$$d_2 = 1$$

$$\mathbf{y} \in \{c_1, c_2, \dots, c_k\}$$

\mathbf{y} is **discrete**

y_i tells you what
class x_i belongs to

Fundamental Supervised ML Problems

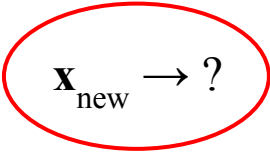
$$\mathbf{x} \in \mathbb{R}^{d_1}, \mathbf{y} \in \mathbb{R}^{d_2}$$

$$\mathbf{x}_1 \rightarrow \mathbf{y}_1$$

$$\mathbf{x}_2 \rightarrow \mathbf{y}_2$$

...

$$\mathbf{x}_n \rightarrow \mathbf{y}_n$$


$$\mathbf{x}_{\text{new}} \rightarrow ?$$

Classification:

$$d_2 = 1$$

$$\mathbf{y} \in \{c_1, c_2, \dots, c_k\}$$

\mathbf{y} is **discrete**

y_i tells you what
class x_i belongs to

Regression:

$$\mathbf{y} \in \mathbb{R}^{d_2}$$

\mathbf{y} is **continuous** (potentially
multi-dimensional)

x_i is being **mapped** to point
 y_i in a different space

Fundamental Supervised ML Problems

$$\mathbf{x} \in \mathbb{R}^{d_1}, \mathbf{y} \in \mathbb{R}^{d_2}$$

$$\mathbf{x}_1 \rightarrow \mathbf{y}_1$$

$$\mathbf{x}_2 \rightarrow \mathbf{y}_2$$

...


$$\mathbf{x}_n \rightarrow \mathbf{y}_n$$

$$\mathbf{x}_{\text{new}} \rightarrow ?$$

Classification:

$$\mathbf{x} = [e_1, e_2, \dots, e_8]^T$$

$$\mathbf{y} = \{\text{"trying to open"}, \text{"relaxing"}\}$$



Fundamental Supervised ML Problems

$$\mathbf{x} \in \mathbb{R}^{d_1}, \mathbf{y} \in \mathbb{R}^{d_2}$$

$$\mathbf{x}_1 \rightarrow \mathbf{y}_1$$

$$\mathbf{x}_2 \rightarrow \mathbf{y}_2$$

...


$$\mathbf{x}_n \rightarrow \mathbf{y}_n$$

$$\mathbf{x}_{\text{new}} \rightarrow ?$$

Regression:


$$\mathbf{x} = [s_1, s_2, \dots, s_5]^T$$

$$\mathbf{y} = [l_a, l_b]^T$$



Nomenclature

$$\mathbf{x} \in \mathbb{R}^{d_1}, \mathbf{y} \in \mathbb{R}^{d_2}$$




\mathbf{x} : “features” / “independent variables”

\mathbf{y} : “labels” / “dependent variables”


“examples” / “training” / “fitting”

“testing” / “predicting”



Neural Networks


- $h_i = f(\sum_j x_j w_{ji}^1)$
- $y_i = \sum_j h_j w_{ji}^2$
- Training the network: finding w_{ij}^k
- Done via backpropagation algorithm



“Fully connected”


Neural Networks

- $h_i = f(\sum_j x_j w_{ji}^1)$
- $y_i = \sum_j h_j w_{ji}^2$
- Training the network: finding w_{ij}^k
- Done via backpropagation algorithm




Neural Networks

- $h_i = f(\sum_j x_j w_{ji}^1)$
- $y_i = \sum_j h_j w_{ji}^2$
- Training the network: finding w_{ij}^k
- Done via backpropagation algorithm



Deep Learning


- Deep Neural Network (DNN): more than one hidden layer
- Even two hidden layers make for incredibly powerful approximators even for highly non-linear functions
- Hidden layers can have different sizes
- Enabled by:
 - better training (backprop)
 - massively parallel computers (GPUs)



Deep Neural Networks




Image pixels
(all of them!)



Deep Neural Networks

- “Network engineering”: setting the number and dimensions of the hidden layers
- Qualitative characteristics:
 - for some problems, they easily outperform anything else
 - very hard to model and analyze quantitatively, to get a sense for what makes them “tick”
 - need large amounts of data to train



What Could I Be Learning?

- Input: **sensor data**
 - Semantic meaning of my sensor data
- Input **sensor data plus action**
 - What will happen next?
 - How “good” is this action?
 - What action should I take?

What Could I Be Learning?

- Input: **sensor data**
 - Semantic meaning of my sensor data
- Input **sensor data plus action**
 - What will happen next?
 - How “good” is this action?
 - What action should I take?

Learning Dynamic Models

- time enters the picture - denoted by subscript t
- system state: $\mathbf{x}_t \in \mathbb{R}^{d1}$
- action / command: $\mathbf{a}_t \in \mathbb{R}^{d2}$
- **model of system dynamics:**
 - $\mathbf{x}_{t+1} = f(\mathbf{x}_t, \mathbf{a}_t)$
 - shorthand: dynamic(s) model, forward model

Learning Dynamic Models

- time enters the picture - denoted by subscript t
- system state: $\mathbf{x}_t \in \mathbb{R}^{d1}$
- action / command: $\mathbf{a}_t \in \mathbb{R}^{d2}$
- **model of system dynamics:**
 - $\mathbf{x}_{t+1} = f(\mathbf{x}_t, \mathbf{a}_t)$
 - shorthand: dynamic(s) model, forward model
- “model-based”: $f(\mathbf{x}_t, \mathbf{a}_t)$ is an analytical function written by the roboticist
- “learning-based”: $f(\mathbf{x}_t, \mathbf{a}_t)$ is learned from data, standard regression problem:

$$(\mathbf{x}_t, \mathbf{a}_t) \rightarrow \mathbf{x}_{t+1}$$

Learning Dynamic Models

- **model of system dynamics:**
 - $\mathbf{x}_{t+1} = f(\mathbf{x}_t, \mathbf{a}_t)$
- “learning-based”: $f(\mathbf{x}_t, \mathbf{a}_t)$ is learned from data, standard regression problem
- why learn dynamic models?
 - can use many non-learning techniques derived over the years that require the existence of one

What Could I Be Learning?

- Input: **sensor data**
 - Semantic meaning of my sensor data
- Input **sensor data plus action**
 - What will happen next?
 - How “good” is this action?
 - What action should I take?

Reinforcement Learning

Conceptually very different from supervised / unsupervised learning.

- system state: $\mathbf{x}_t \subset \mathbb{R}^{d1}$ - **known**
- action / command: $\mathbf{a}_t \subset \mathbb{R}^{d2}$ - **controllable**
- dynamics / forward model:
 - deterministic: $\mathbf{x}_{t+1} = f(\mathbf{x}_t, \mathbf{a}_t)$
 - stochastic: $P(\mathbf{x}_{t+1} | \mathbf{x}_t, \mathbf{a}_t) = f(\mathbf{x}_t, \mathbf{a}_t)$
 - could be **known** (model-based RL) or **unknown** (model-free RL)
- reward: $r(\mathbf{x}_t, \mathbf{a}_t)$ - known
- goal: learn a **policy** that maximizes **expected reward**
 - deterministic: $\mathbf{a}_t = \Pi(\mathbf{x}_t)$
 - stochastic: $P(\mathbf{a}_t | \mathbf{x}_t) = \Pi(\mathbf{a}_t, \mathbf{x}_t)$

Reinforcement Learning

Conceptually very different from supervised / unsupervised learning.

- reward: $r(\mathbf{x}_t, \mathbf{a}_t)$ - known
- goal: learn a **policy** that maximizes **expected reward**
 - deterministic: $\mathbf{a}_t = \Pi(\mathbf{x}_t)$
 - stochastic: $P(\mathbf{a}_t | \mathbf{x}_t) = \Pi(\mathbf{a}_t, \mathbf{x}_t)$
- robot is never told “how to do something”, only if what it did is good or bad
- many algorithms:
 - Q-learning
 - policy gradient
 - ...

How Can I Apply ML to My Robotics Problem?

Key questions:

- **What kind of a problem is it?** Clustering, dimensionality reduction, regression, classification, policy.

How Can I Apply ML to My Robotics Problem?

Key questions:

- **What kind of a problem is it?** Clustering, dimensionality reduction.
- **What are the features?** How do I get a feature vector $\mathbf{x} \in \mathbb{R}^{d1}$ from my data?

How Can I Apply ML to My Robotics Problem?

Key questions:

- **What kind of a problem is it?** Regression, classification.
- **What are the features?** How do I get a feature vector $\mathbf{x} \in \mathbb{R}^{d1}$ from my data?
- **What are the labels?** How do I express what I care about as $\mathbf{y} \in \mathbb{R}^{d2}$?
- **Where does my training data come from?** It needs to be labeled with “ground truth”.

How Can I Apply ML to My Robotics Problem?

Key questions:

- **What kind of a problem is it?** Policy
- **What is my state space?** It needs to be known for training **and** testing.
- **What is my action space?** It needs to contain things I can control.
- **What is my reward function?** I need to have access to a reward signal, either by computing it, or from an oracle.

How Can I Apply ML to My Robotics Problem?

For all problems:

- **What algorithm/model is being used?**
 - “fit” my data well
 - run in reasonable time on my problem size
 - have low training loss
- **How are the results tested?**
 - what is the performance on the testing data?
 - what is the relationship between training and testing data?
 - *is testing data predictive of deployment?*

Course Structure

- Components (roughly 7-10 lectures each):
 - Part 1: Dimensionality reduction, classification and regression
 - Part 2: Deep Learning and learning in Computer Vision
 - Part 3: Reinforcement learning
- Lectures
 - theoretical concept presentation (in moderate depth)
 - discussion of papers in Robotics using these concepts

Course Structure

- 3 Programming Projects: one for each part
 - Each project will make up ~15% of the grade.
 - All projects: Python, Ubuntu 16.04, ROS, SVN. **Must be comfortable with all of these.**
 - There will be no skeleton code. You will write ROS nodes starting from blank files.
 - You will get a budget of 5 late days on projects, to use how you see fit.
- Written exams
 - Midterm and final: each will make up ~20% of the grade.
 - Pencil-and-paper questions: conceptual, algorithmic, math-y

Additional details for course structure, grading, collaboration policy, late submissions, etc. **can be found on the Courseworks2 course webpage.**